COGNITRA

Structure Before Synthesis

A Controlled Intelligence Architecture for Decision-Making

Briefs are rendered from validated records — not from PDFs.

HOW MOST AI SYSTEMS WORK

```
Raw Documents → Black-Box
Model → Narrative Output
```

RISKS

- Hallucinations
- No validation
- No governance
- No traceability
- Non-repeatable outputs

COGNITRA SEPARATES FACT FROM STRATEGY

1 **Structured Extraction (LLM)**
PDF → Strict JSON Schema
Factual fields only

↓

2 **Deterministic Postprocessing (Code)**
Priority · Confidence · Region Mapping · Macro-Themes
Rule-based. Reproducible. Auditable.

↓

3 **Human Review Gate**
Approved records only

↓

4 **Executive Brief Rendering**
Generated from **approved JSON only**
Never from raw PDFs

WHAT THIS ENABLES

- Traceable Insights
- Controlled Interpretation Boundaries
- Reproducible Executive Outputs
- Governance-Ready AI
- Reusable Intelligence Layer
- Cross-Record Signal Detection

ARCHITECTURE PRINCIPLES

- Structure before synthesis
- Human approval gate
- Cost-aware model usage

- Deterministic validation
- Controlled interpretation
- Governance-ready outputs

Cognitra builds a validated intelligence layer — then synthesizes.